

Introduction to Operating Systems

M1 MOSIG – Operating System Design

Renaud Lachaize

Acknowledgments

- Many ideas and slides in these lectures were inspired by or even borrowed from the work of others:
 - Arnaud Legrand, Noël De Palma, Sacha Krakowiak
 - David Mazières (Stanford)
 - Many slides directly adapted from those of the CS140 class
 - Remzi and Andrea Arpaci-Dusseau (U. Wisconsin)
 - Textbooks (Silberschatz et al., Tanenbaum)

Course goals

- **Introduce you to operating system concepts**
 - Hard to use a computer without interacting with the OS
 - Understanding the OS makes you a better (more effective) programmer
- **Cover important system concepts in general**
 - Caching, concurrency, memory management, I/O, protection, ...
- Teach you to deal with larger software systems
- Prepare you to take other classes related to OS concepts
 - M1 Principles of computer networks, M1/M2 Distributed systems, M2 Parallel systems, M2 Advanced OS, M2 cloud infrastructure, ...

Outline

- What is an operating system?
- Some history
- Abstractions: processes and address spaces
- Protection and resource management

What is an operating system?

- *An operating system (OS) is a (software) layer between the hardware and the applications*
- Two key roles: **virtualization** and **resource management**
- **Virtualization**
 - *The OS makes it easier to write and run programs on a machine*
 - Hides the low-level interface of the hardware and replaces it with higher-level abstractions
 - Hides the physical limitations of a machine and the differences between machines (size of the main memory, number of processors)
 - Hides the sharing of resources between applications/users
 - Thus, we sometimes refer to the OS as a “*virtual machine*”

What is an operating system? (continued)

- **Resource management**
 - The OS is in charge of managing the resources of a computer system
 - **Physical resources:** memory, processor, devices, ...
 - **Logical resources:** programs, data, communications, ...
 - Goals: allow the applications to run safely / securely / efficiently / fairly ... despite the fact that they run concurrently
 - Encompasses several dimensions, including:
allocation, sharing and protection
 - Consists in a combination of ***mechanisms*** and ***policies***

OS Design goals and trade-offs

- Provide **useful abstractions** to improve programmer/administrator/user productivity
- Provide high **performance**
 - Leverage the power/capacity of the hardware
 - Minimize the (time and space) overhead of the OS features
- Provide **protection**
 - Between applications
 - Between applications and OS
 - Between users
- Provide a high degree of **reliability**
- Take care of other aspects such as predictability, energy-efficiency, mobility, ...

OS Interfaces

- An operating system typically exports **two kinds of interfaces**
 - A command/user interface
 - A programmatic interface
- **Command/user interface**
 - Designed for human users
 - Various forms: textual or graphical
 - Composed of a set of commands
 - Textual example (Unix shell): `rm myfile.txt`
 - Graphical example (most systems): drag the myfile.txt icon into the trash bin.

OS Interfaces (continued)

- **Programmatic interface**
 - This interface is used/called from application programs running on the system
 - Including the programs implementing command/user interfaces
 - Composed of a set of procedures/functions
 - Libraries
 - System calls (more details later)
 - Defined both:
 - At the source code level: *Application Programming Interface (API)*
 - At the machine code level: *Application Binary Interface (ABI)*

Some of the topics that we will study
during the semester

- **How does the OS virtualize and manage resources?**
 - What are the required mechanisms and policies?
 - What kind of support is required from the hardware?
 - How can these goals be achieved efficiently?
 - We will consider several resources : CPU, main memory, input/output (I/O) devices (e.g., storage devices)
- **How to build concurrent programs?**
 - How to program applications with several “tasks”?
 - How to coordinate these tasks and let them share data?
 - How to make such programs correct and efficient?
 - What kind of support is needed from the OS and the hardware to achieve this goals?

Outline

- What is an operating system?
- **Some history**
- Abstractions : processes and address spaces
- Protection and resource management

Some History

(1) Early operating systems

- In the beginning, **the OS did not do much**
- Essentially, a set of libraries of commonly-used functions (e.g., low-level code for I/O devices)
- **Running one program at a time**
- Possibly involving a human operator (e.g., for deciding in what order to run the jobs)
- **Assumed no bad users or programs**
- **Problem: poor utilization**
 - ... of hardware (e.g., CPU idle while waiting for I/O completion)
 - ... of human user (must wait for each program to finish)

Some History

(2) Beyond libraries: Protection

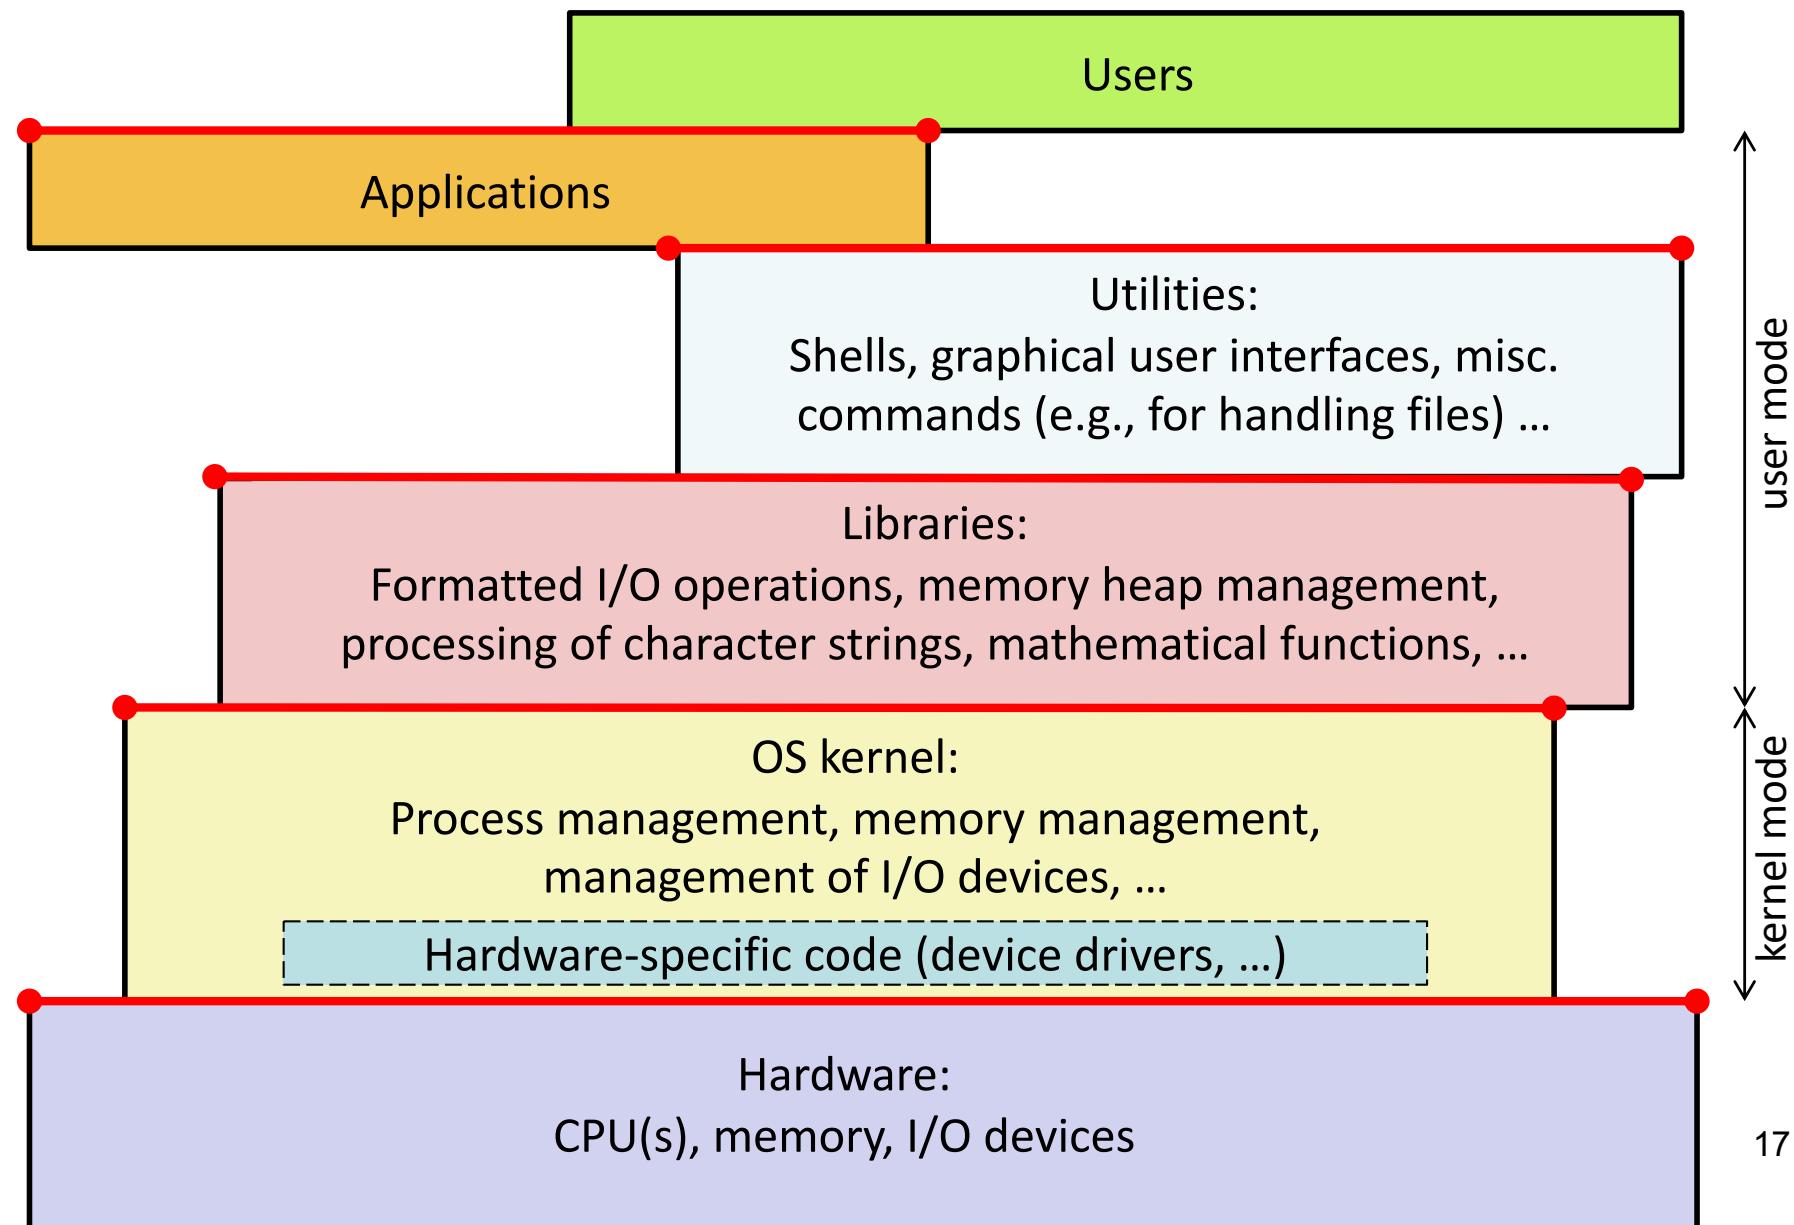
- Realization that the code of the OS plays a central role
- **A user/application should not be able to make the whole system fail or to perform unauthorized operations**
 - E.g., issue arbitrary write requests to a storage device
- Idea: Modification of the OS interface
 - Old interface: provide applications with library procedures allowing direct access to critical operations
 - New interface: **force application to delegate critical operations, using a hardware mechanism** that transfers control to a more privileged execution mode
 - Such an interface is called a “system call” or “syscall” (more details later)

Some History

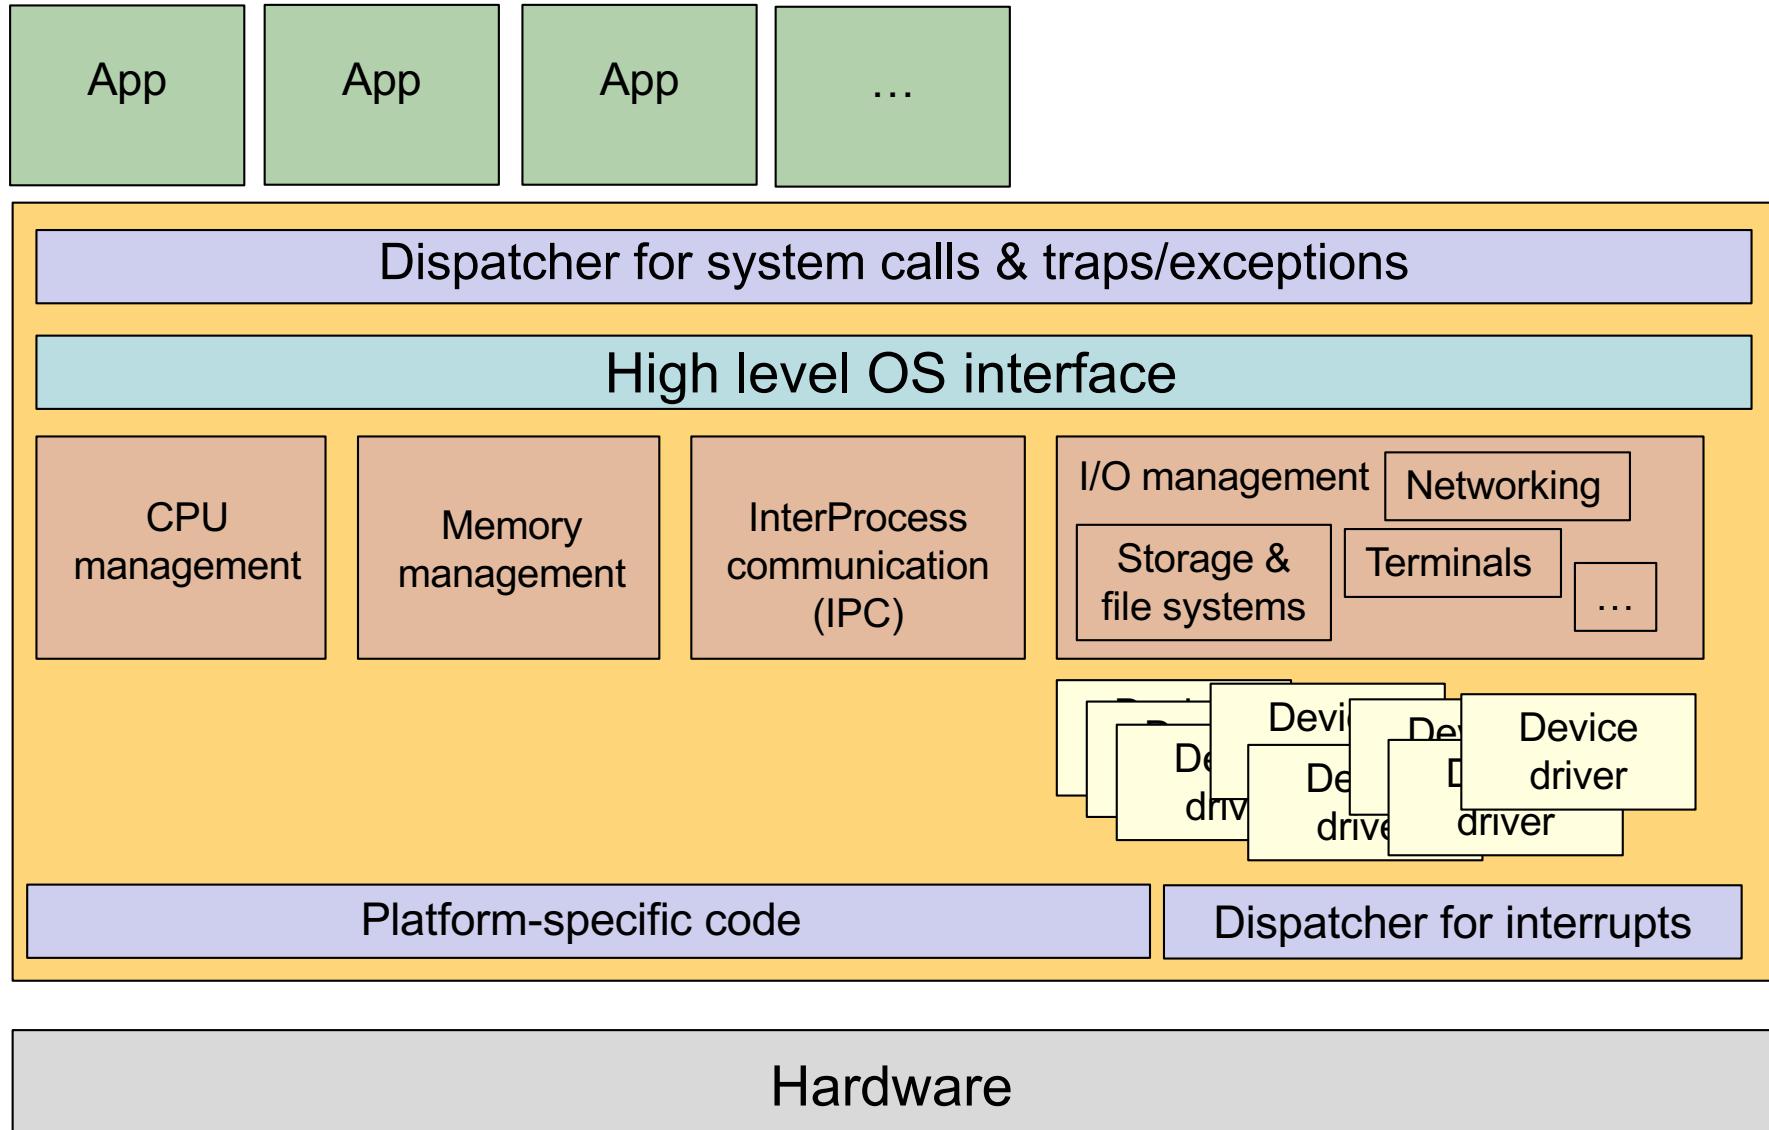
(3) Multiprogramming / Multitasking

- Idea: **improve machine resource utilization by running several programs concurrently**
 - When a program blocks (e.g., waiting for input from the disk / the network / the user), run another program
- **Problems:** what can an ill-behaved application do ?
 - Never relinquish the CPU (infinite loop)
 - Access the memory of another application
- **The OS provides mechanisms to address these problems**
 - Preemption: take CPU away from a looping application
 - Memory protection: prevent an application from accessing another application's memory

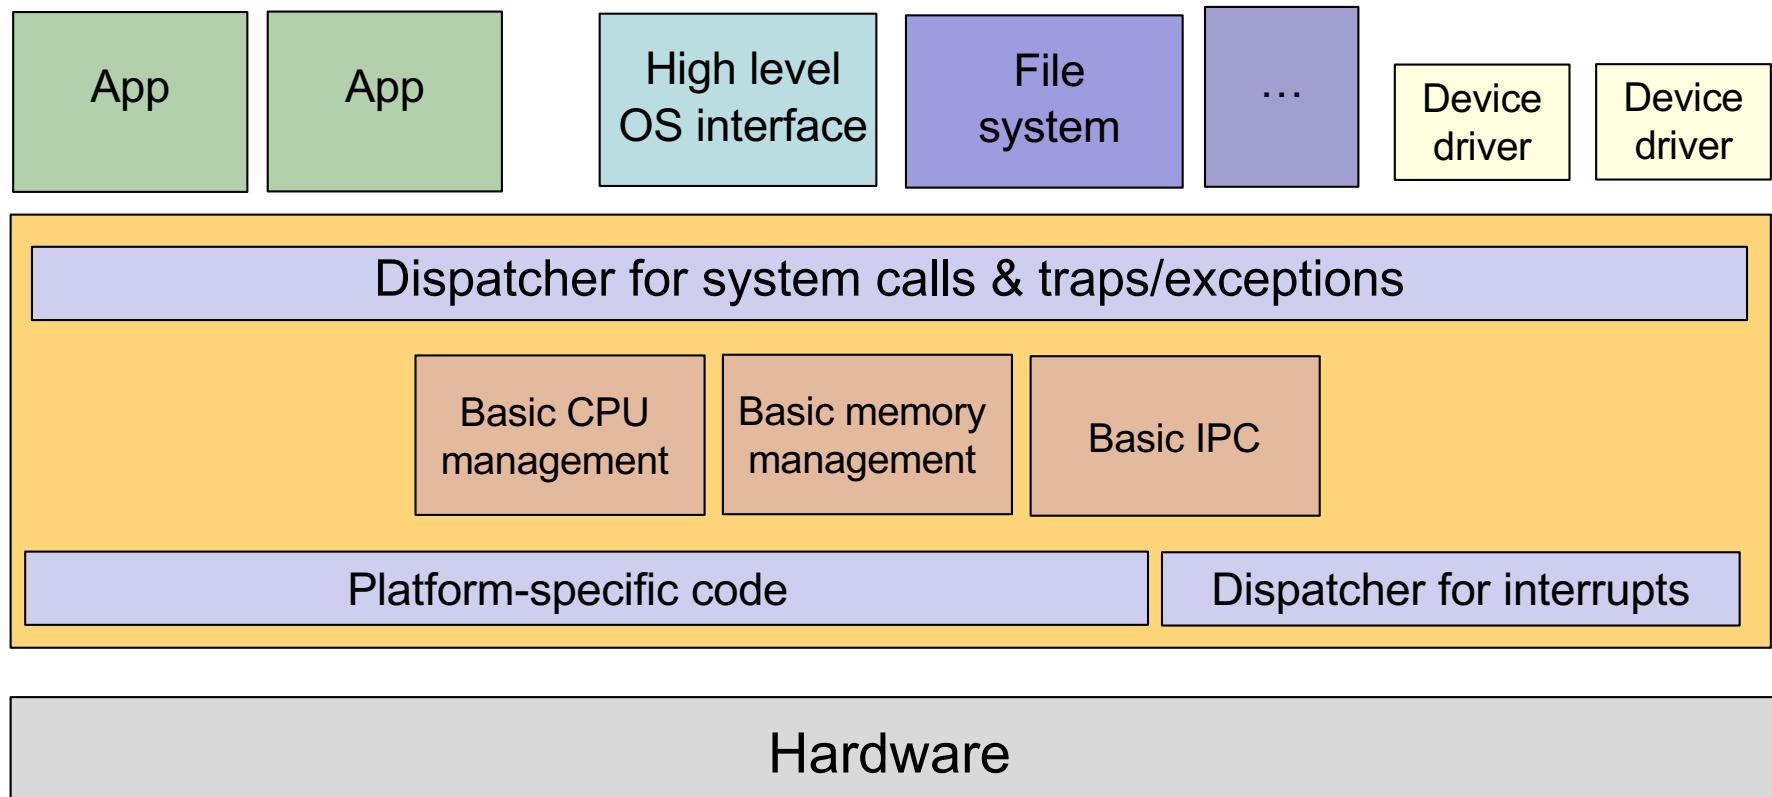
Typical structure of an operating system



Monolithic kernel design (e.g., Linux)



Microkernel design (e.g., L4)



Outline

- What is an operating system?
- Some history
- **Abstractions: processes and address spaces**
- Protection and resource management

A key OS abstraction: the Process

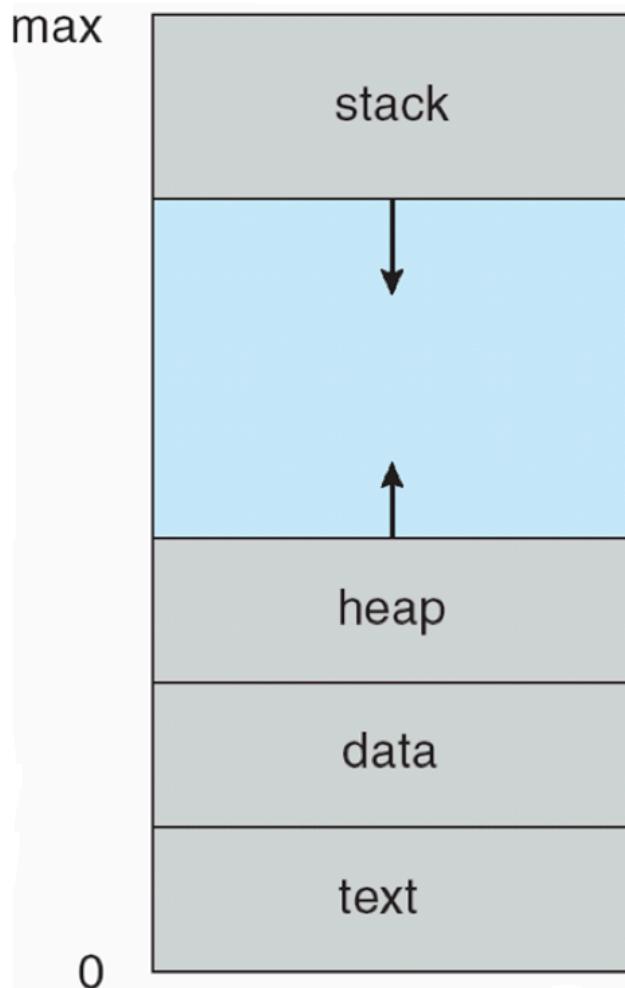
- **A process is an abstraction corresponding to a running instance of a program**
- Its main role consists in **virtualizing the CPU**
 - Although there are just a few physical CPUs (or even just one), the OS can provide the illusion of a nearly-endless supply of logical CPUs (one per process)
 - It also allows the OS to capture the state and control the execution of a running program, which are key mechanisms for resource management

A key OS abstraction: the Process

- A process mainly consists in:
 - An ***execution context*** (a.k.a. an *execution flow*, or a *control flow*):
 - A current machine state: a set of current values for the CPU registers, including the program counter (PC) and the stack pointer (SP)
 - An execution stack
 - A ***memory space*** (a.k.a. an *address space*)
 - A ***logical state*** (is it currently running? If not, why?)
 - Some other information, required by the OS

Process address space

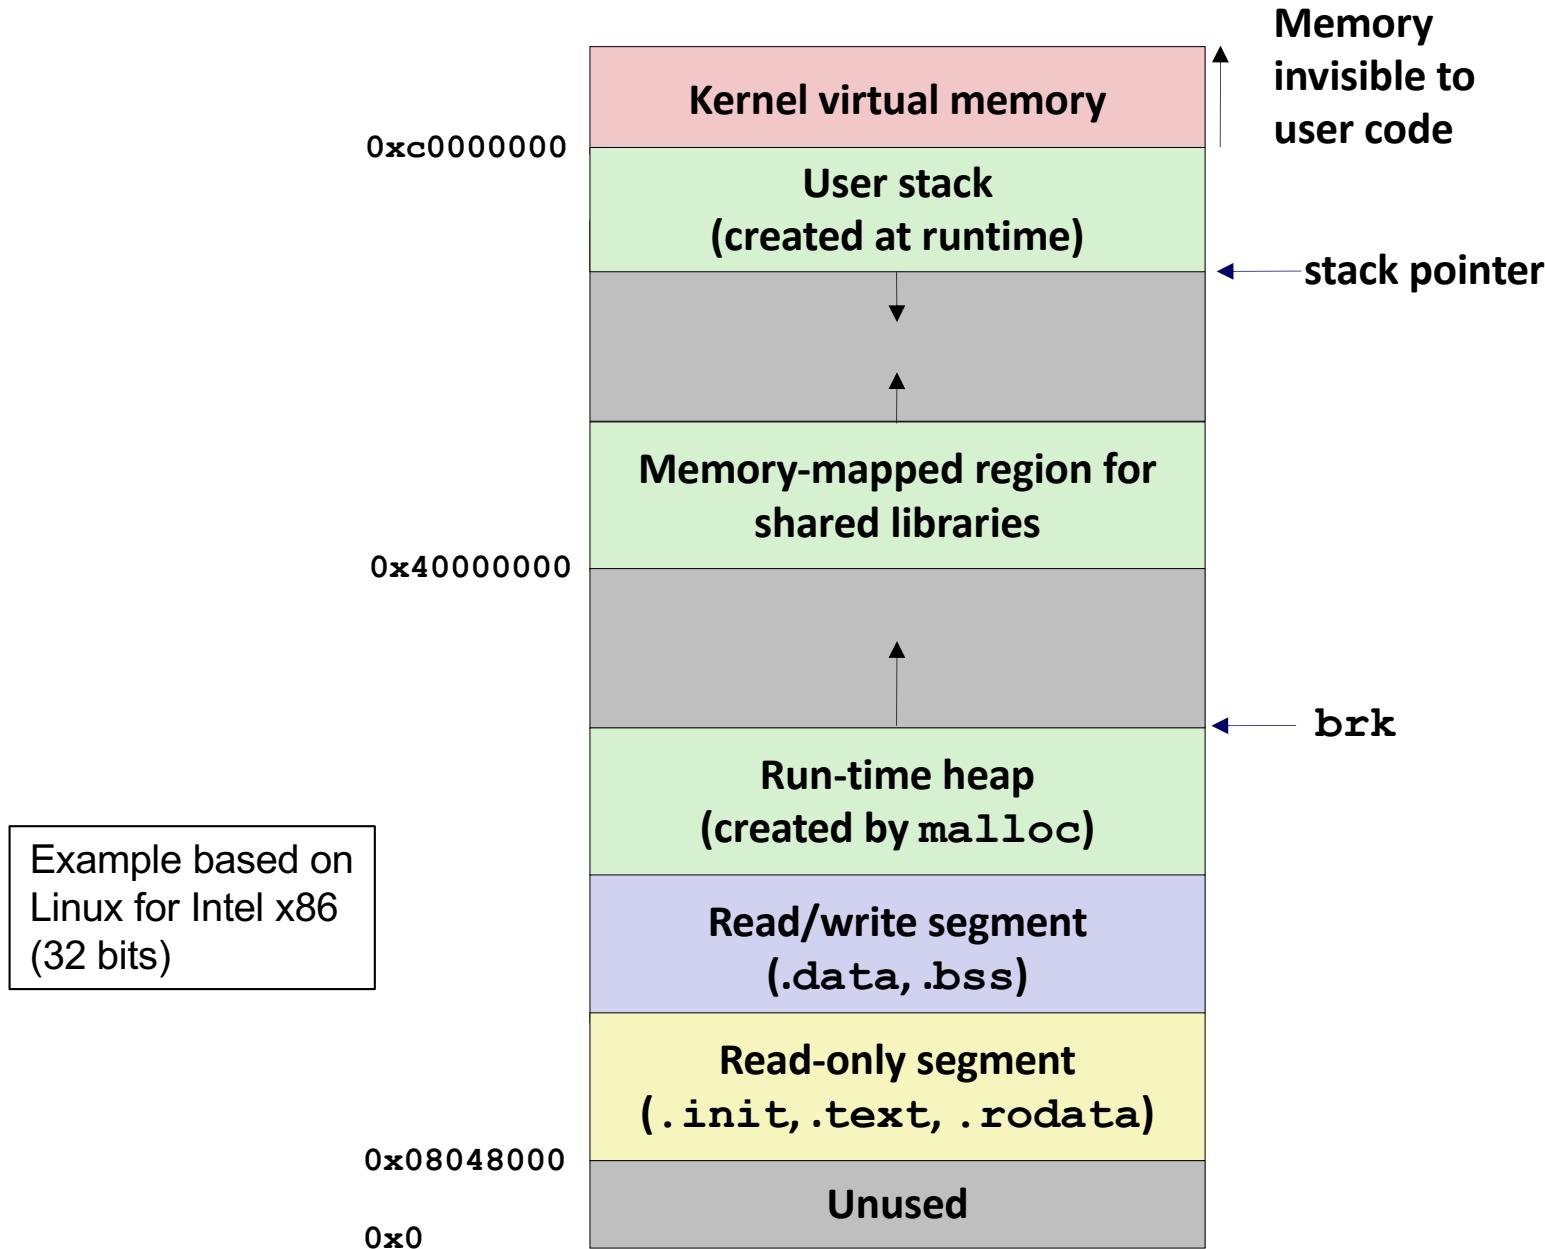
A simplified view



Picture from: Silberschatz et al., *Operating systems concepts* (8th edition)

Process address space

A more detailed view



Outline

- What is an operating system?
- Some history
- Abstractions : processes and address spaces
- **Protection and resource management**

Some key techniques for protection

- Overall goal: **prevent bad processes from impacting the OS or other processes**
- **Preemption**
 - **Give a resource to a process and take it away if needed** for something else
 - Example: CPU preemption
- **Interposition**
 - **Place OS between application and resources** (e.g., an I/O device, or a piece of information stored in memory)
 - OS tracks the resources that the application is allowed to use
 - On every access request, check that the access is legal
 - Example: System calls

Some key techniques for protection (continued)

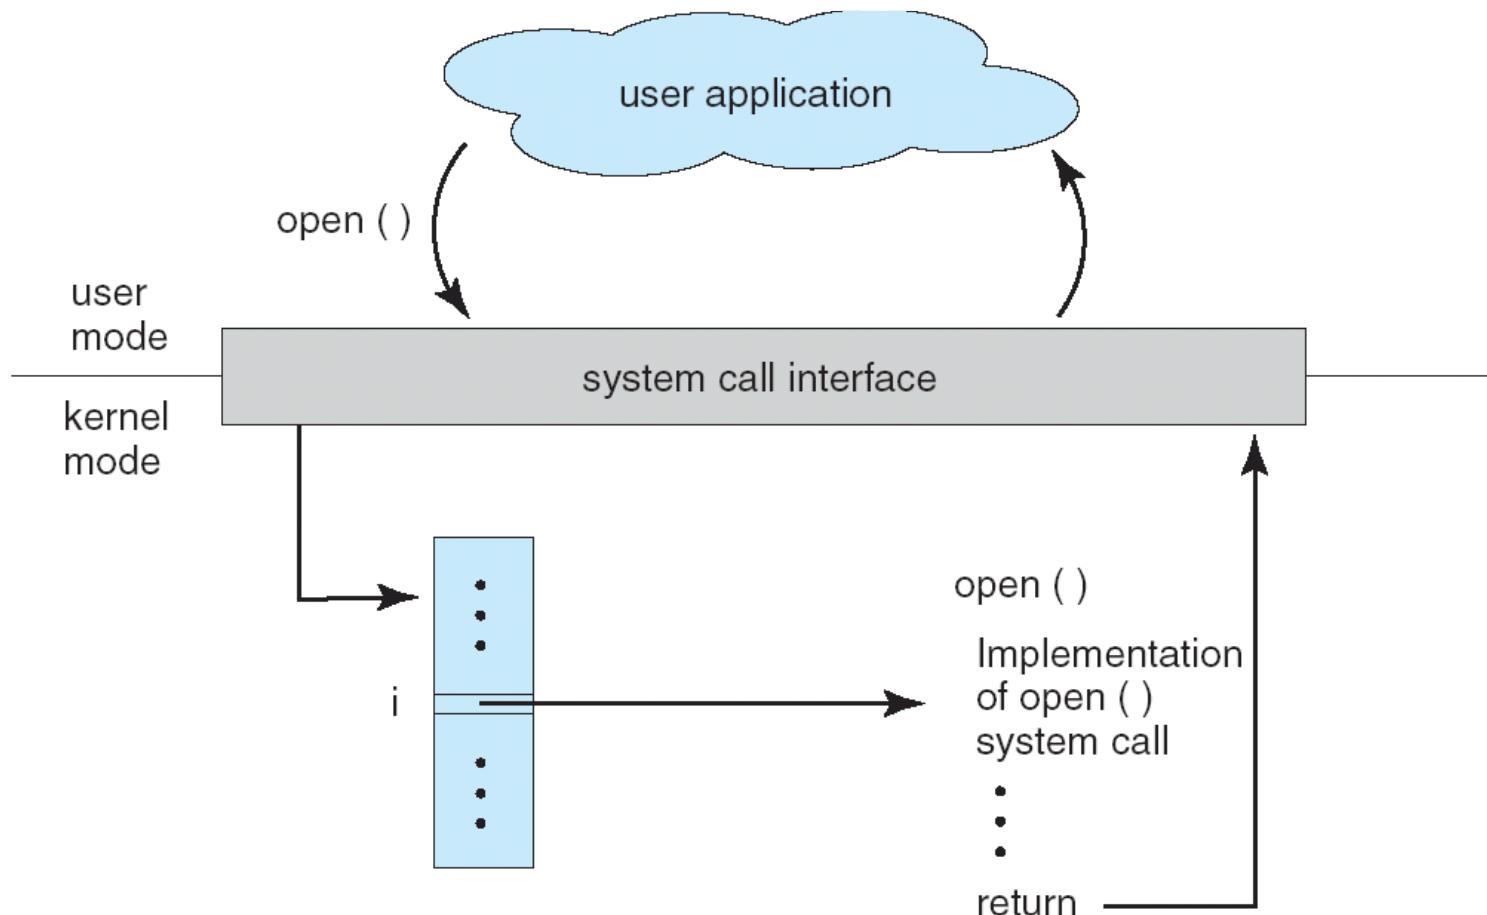
- CPU execution modes
 - CPUs provide **2 execution modes**:
 - Privileged (a.k.a. supervisor mode, or kernel-level mode)
 - Unprivileged (a.k.a. user mode, or user-level mode)
 - OS kernel code runs in **privileged mode**
 - Application code runs in **unprivileged mode**
 - Protection-related code (resp. data) must only be executed (resp. accessed) in privileged mode
 - Enforced by hardware (details later)
 - A system call is the only way to switch from unprivileged mode to privileged mode

System calls

- Applications (i.e., user-level code) can **invoke kernel services** through the system call mechanism
 - **Using a special hardware instruction** that triggers a trap into kernel-mode
 - ... and transfers control to a trap handler
 - ... which dispatches to one of a few hundred syscall handlers

System calls (continued)

- Illustration with the `open` system call (to open a file)



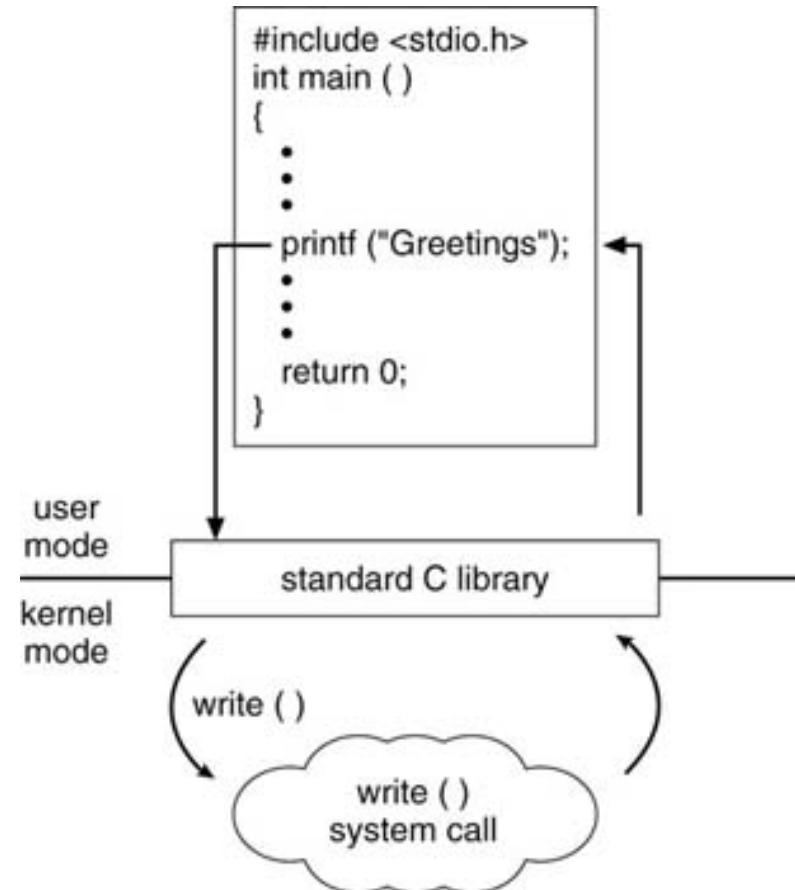
Picture from: Silberschatz et al., *Operating systems concepts* (8th edition)

System calls (continued)

- Goal: **perform things that an application is not allowed to do in unprivileged mode**
 - Like a library call, but into more privileged code
- The kernel supplies a **well-defined system-call interface**
 - Applications set up syscall arguments and trap to kernel
 - Kernel checks if operation is allowed, performs operation and returns results (transfers control back to application)
 - Many higher-level library functions are built on the syscall interface
 - Example (Unix) : functions such as `printf` and `scanf` are implemented as user-level library code that calls the kernel using system calls such as `read` and `write`

System call example

- The standard C library (libc) is implemented in terms of syscalls
 - **printf** (in libc) has same privilege as application
 - **printf** calls **write**, which can access low-level resources such as the console/screen and files



CPU preemption

- **Protection mechanism to prevent a process from monopolizing the CPU**
 - Allows the kernel to take back control of the CPU after a maximum time interval
 - Relies on the processor interrupt mechanism and on a timer device
- **The kernel programs the timer to send periodic interrupts (e.g., every 10 ms)**
 - Device configuration is only allowed in privileged mode
 - User code cannot re-program the timer

CPU preemption (continued)

- The kernel configures the processor to define a **timer interrupt handler**
 - This handler is run in privileged mode
 - In this way, each periodic timer interrupt will trigger the execution of some kernel-defined code
 - This kernel code can decide to keep the current process running or to give the CPU to another one
 - Note : interrupt entry points cannot be defined/modified by user-level code
 - Thus, there is no way for user code to hijack the interrupt handler
- Result: a process cannot monopolize the CPU with an infinite loop
 - At worst, it may get $1/N$ of CPU time if there are N CPU-hungry processes

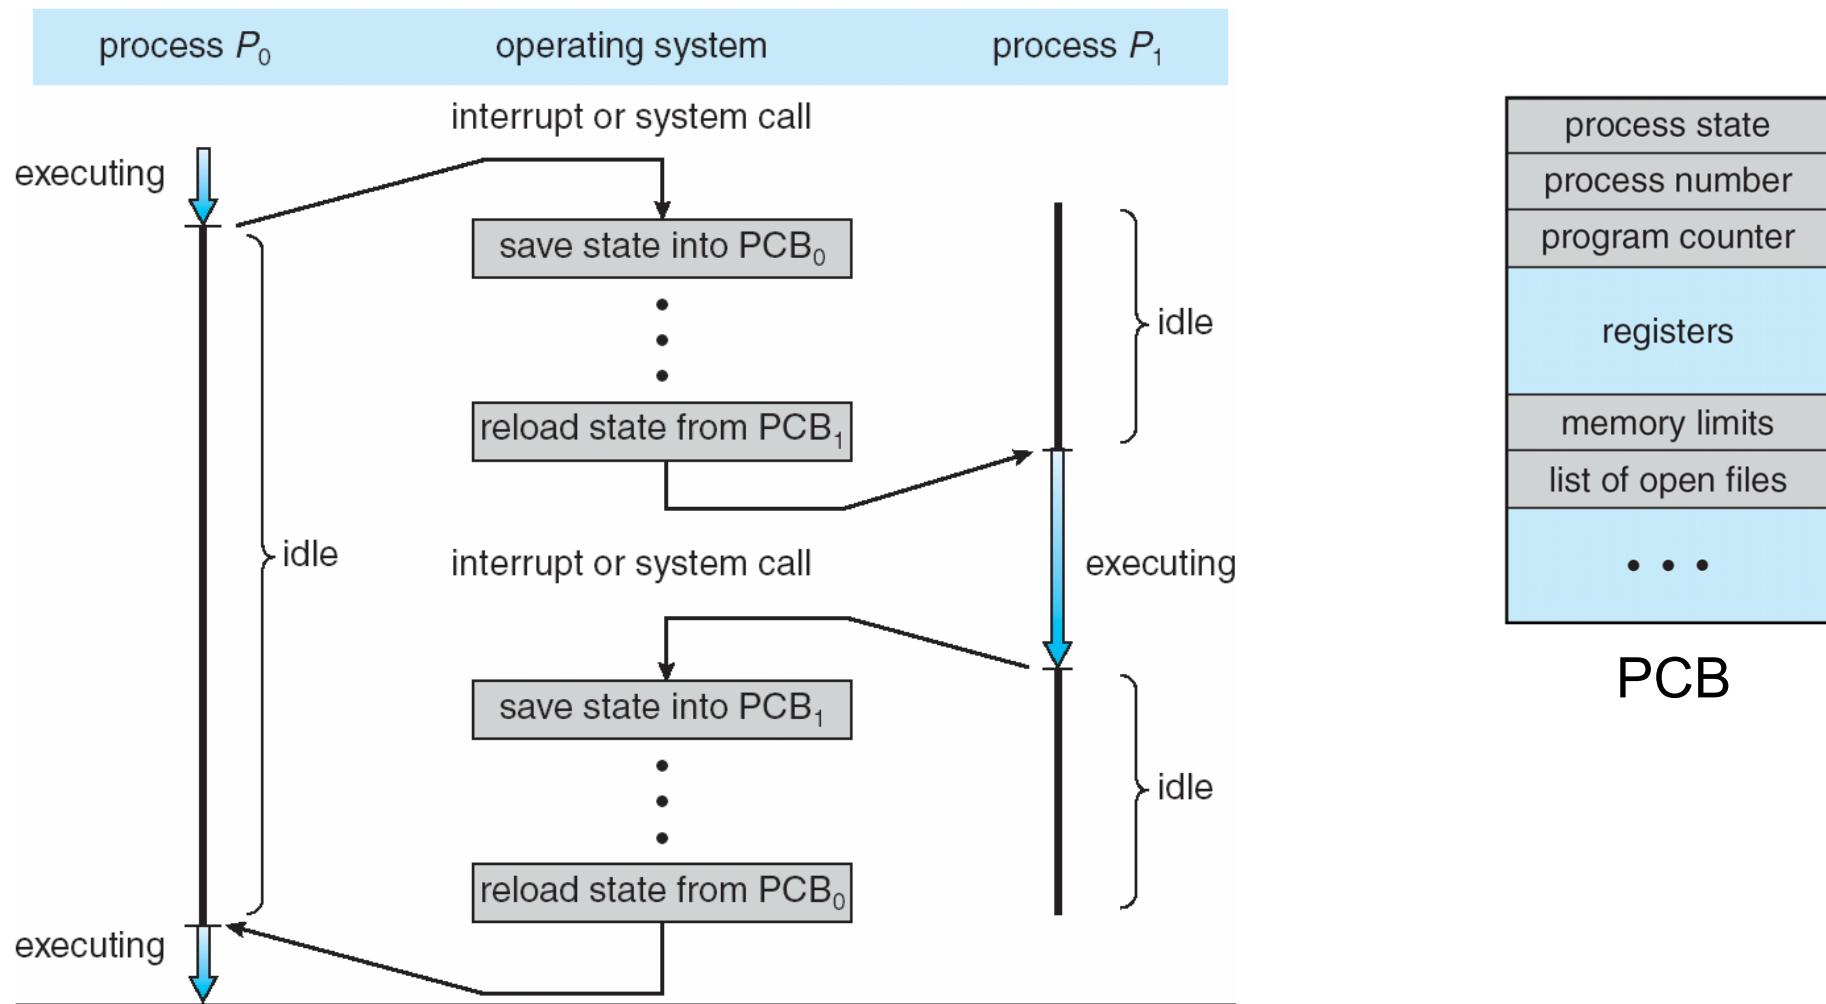
CPU scheduling

- **The *scheduler* is a component of the OS, in charge of deciding which process should run on the CPU (1 decision per CPU)**
- When is the scheduler invoked?
 - Periodically, for each timer interrupt
 - Punctually, in reaction to some syscalls:
 - Process termination (exit)
 - Process explicitly releasing the CPU (yield, sleep, ...)
 - Process requesting a blocking action
 - Creation of a new process with a higher priority
 - ...
 - Punctually, in reaction to some interrupts
 - E.g., a device notification for available data

CPU scheduling (continued)

- **What does the scheduler do upon invocation?**
 - **Make decision on the process P2 that should obtain the CPU**, based on:
 - The list of processes that are ready to run
 - ... and a given scheduling policy
 - **Save execution context of “outgoing” process P1**
 - (Except if this process is terminated)
 - This allows resuming the execution of P1 later on
 - **Inject /restore the execution context of P2 on the CPU**
- This sequence of steps is called a “**context switch**”
 - Note that, just after the switch, P2 runs in kernel mode and must eventually switch back to user mode. This will happen via a return-from-interrupt or a return-from-syscall instruction.

Context switch

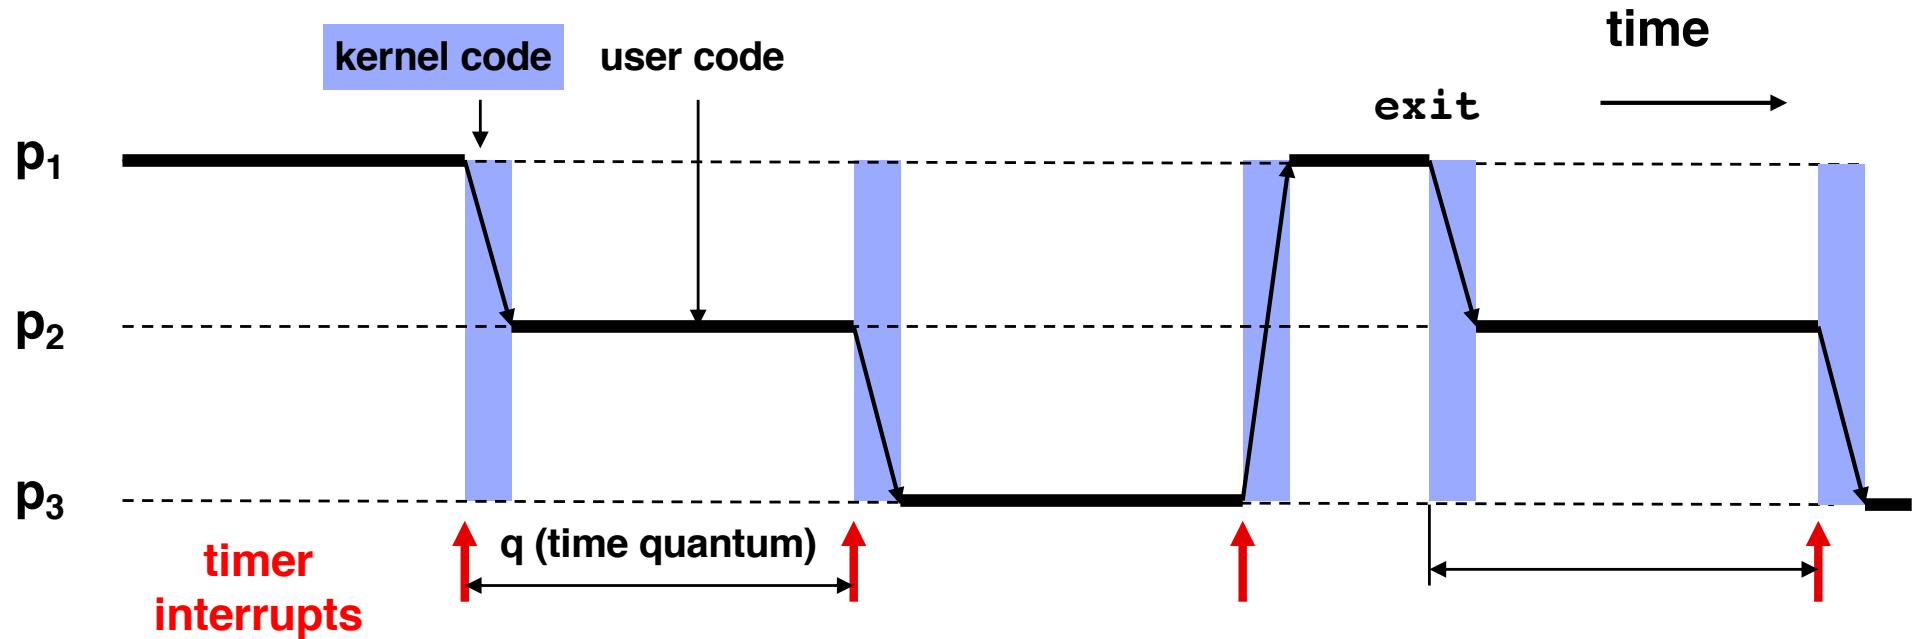


Pictures from: Silberschatz et al., *Operating systems concepts* (8th edition)

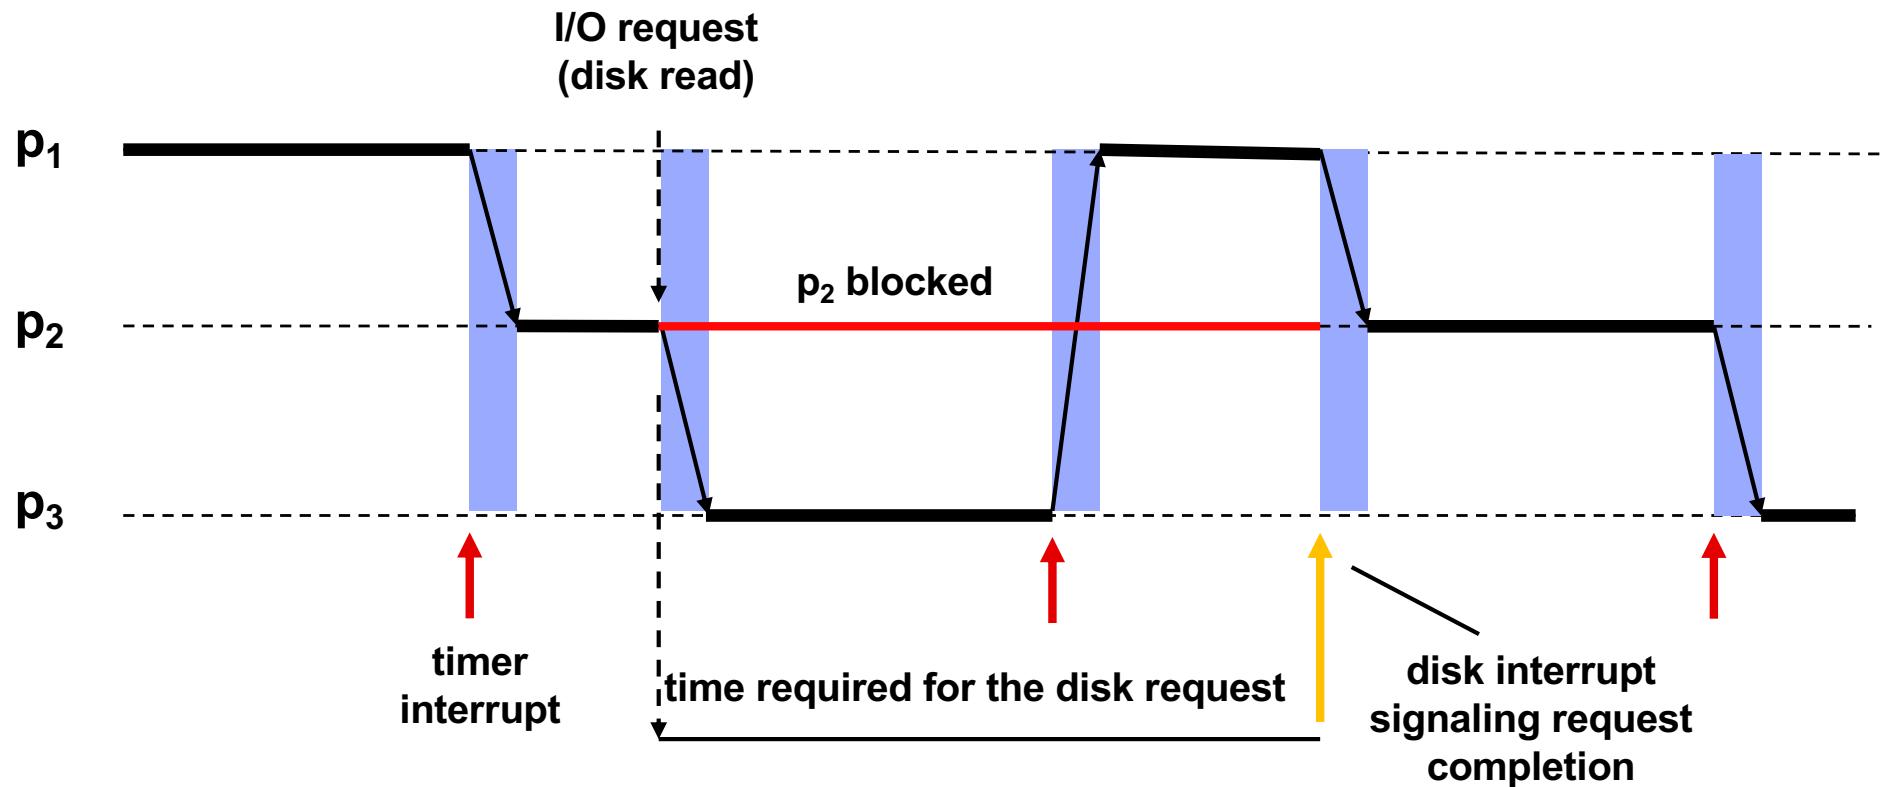
Context switch

- Notes:
 - Implementation details are very machine (processor) dependent, but the general principle is the same
 - A context switch has **a non-negligible cost** and should not happen too often
 - **Warning: Do not confuse**
 - **Context-switch** (transition between two execution contexts)
 - **Mode switch** (transition between user and kernel mode, in the same execution context)

CPU scheduling examples



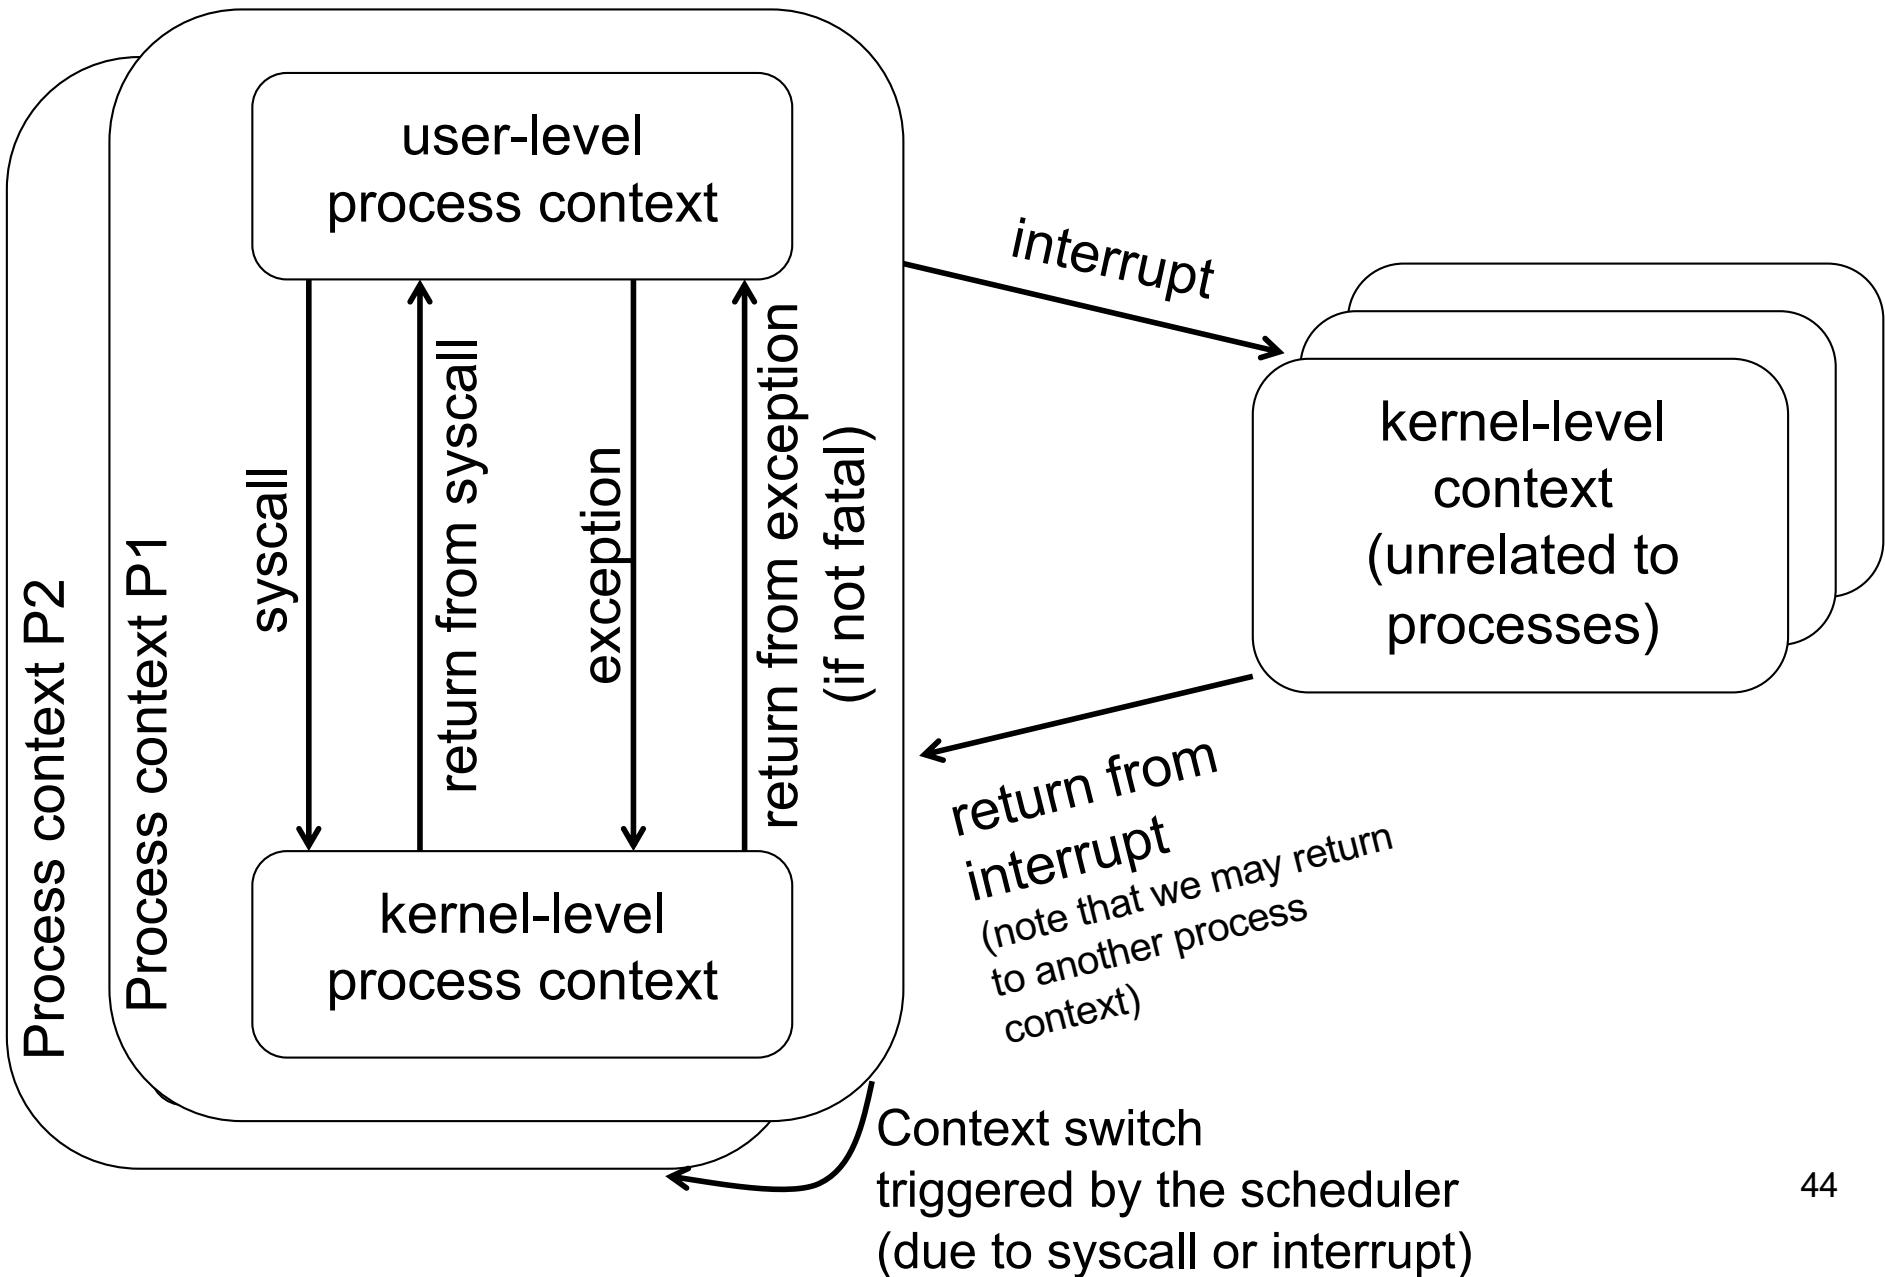
CPU scheduling examples (continued)



Different system “contexts”

- A system can typically be in one of several contexts:
 - User-level process context
 - Running application code or library code called by application
 - Kernel process context
 - Running kernel code on behalf of a particular process
 - Typically, performing a system call
 - Also, exception handling (numeric exceptions, memory faults, ...)
 - Kernel code not associated with a process
 - Timer interrupt handler
 - Device interrupt handlers
 - (Also some handlers for special kinds of interrupts triggered by software)

Different system “contexts” (2)



Memory virtualization and protection

- **The OS must protect the memory space of a process from the actions of other processes**
- Definitions
 - **Address space**: all memory locations that a program can name
 - **Virtual address**: an address in a process address space
 - **Physical address**: an address in real memory
 - **Address translation** : map virtual address to physical address
- **Translation performed for each executed instruction that issues a memory access**
 - Modern CPUs do this in hardware for speed
- Idea: if you cannot name it, you cannot touch it
 - Ensure that the translations of a process do not include memory areas of other processes

Memory virtualization and protection (continued)

- **CPU allows kernel-only virtual addresses**
 - The kernel is typically part of all address spaces, e.g., to handle a system call in the same address space
 - But the OS must ensure that applications cannot touch kernel memory
- **CPU allows disabled virtual addresses**
 - Helps catching and halting buggy program that makes wild accesses
 - Makes virtual memory seem bigger than physical (e.g., bring a page in from disk only when accessed)
- **CPU allows read-only virtual addresses**
 - E.g., allows sharing of code pages between processes
- **CPU allows disabling execution of virtual addresses**
 - Makes certain (code injection) security attacks harder

Summary

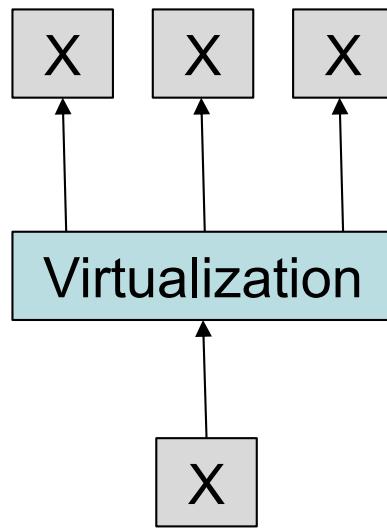
- The main roles of an OS are **virtualization** and **resource management**
- Protection is a fundamental concern
- Some **key abstractions**
 - Processes
 - Virtual address spaces
- Some **key mechanisms (hardware-assisted)**
 - Privileged/unprivileged execution modes
 - System calls and traps
 - CPU preemption (relying on processor interrupts)
 - Memory translation (implementation will be studied in next lectures)

Appendix

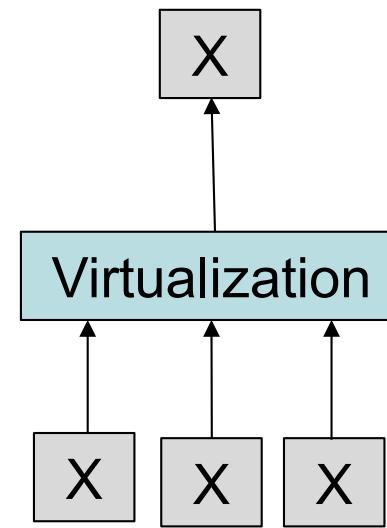
Main techniques for virtualization (1/5)

- In order to virtualize the resources of a machine, operating systems rely on **a combination of 3 main techniques**:
 - Multiplexing (in space and/or in time)
 - Aggregation
 - Emulation
- Note: these techniques are sometimes also used within some hardware devices.

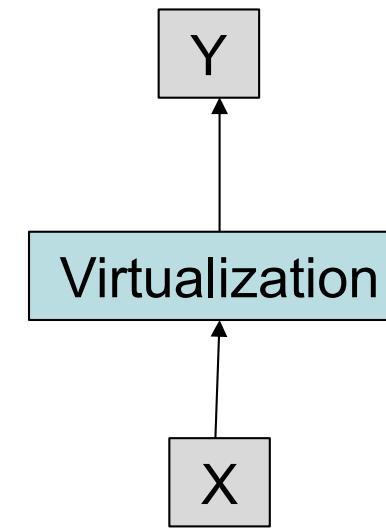
Main techniques for virtualization (2/5)



Multiplexing



Aggregation



Emulation

Main techniques for virtualization (3/5)

- Multiplexing:
 - Exposes a resource among multiple virtual entities
 - Two types of multiplexing: in space and in time
 - Examples:
 - CPUs (in time)
 - Memory (in space and possibly also in time with swapping)
 - I/O devices (in time)

Main techniques for virtualization (4/5)

- Aggregation:
 - The opposite of multiplexing
 - Takes multiple resources and makes them appear as a single abstraction
 - Examples:
 - Memory controller with several DIMMS (hardware)
 - RAID (hardware or software)

Main techniques for virtualization (5/5)

- Emulation
 - Expose (using a level of indirection in software) a virtual resource that is not provided by the underlying machine
 - Examples
 - Sockets and files provide higher-level abstractions above hardware devices
 - Binary translation (compatibility layer) [hardware level]