

Threads


M1 MOSIG – Operating System Design

Renaud Lachaize

Acknowledgments

- Many ideas and slides in these lectures were inspired by or even borrowed from the work of others:
 - Arnaud Legrand, Noël De Palma, Sacha Krakowiak
 - David Mazières (Stanford)
 - **(many slides/figures directly adapted from those of the CS140 class)**
 - Remzi and Andrea Arpaci-Dusseau (U. Wisconsin)
 - Randall Bryant, David O'Hallaron, Gregory Kesden, Markus Püschel (Carnegie Mellon University)
 - Textbook: Computer Systems: A Programmer's Perspective (2nd Edition) a.k.a. "CSAPP"
 - CS 15-213/18-243 classes (many slides/figures directly adapted from these classes)
 - Textbooks (Silberschatz et al., Tanenbaum)

Threads



- **A thread is a schedulable execution context**
 - Program counter, stack, registers ...
- By default, a process uses only one thread
- But it is also possible to have multi-threaded processes
 - Multiple threads running in the same memory address space

Why threads?

- **Most popular abstraction for concurrency**
 - All threads in a process share memory and file descriptors
 - A lighter-weight abstraction for communication than inter-process communication mechanisms
 - Lower resource consumption: a process context requires more resources (memory, initialization and context switching time) than a thread context
- **Allows a process to use multiple CPUs (parallel execution)**
- **Allows a program to overlap I/O and computation**
 - Do not block the whole process when only a part of it should be blocked
 - E.g., a threaded Web server can service several clients simultaneously

```
for(;;) {  
    fd = accept_client();  
    thread_create(service_client, fd)  
}
```

Thread package (pseudo) API

- `tid thread_create (void (*fn) (void *), void *arg);`
 - Create a new thread, run `fn` with `arg`
- `void thread_exit();`
 - Destroy current thread
- `void thread_join(tid thread)`
 - Wait for thread `thread` to exit

- And also lots of support for synchronization (see next lectures)
- **Design choices** (details on next slides):
 - A given thread package can provide either preemptive or non-preemptive (a.k.a. cooperative) threads
 - Kernel-level threads versus user-level threads

Preemptive vs. cooperative threads

- **Preemptive threads**
 - A **thread can be preempted at any time** in order to allocate the CPU to another execution context, e.g., another thread (from the same process) or another process
 - Rely on time multiplexing, thanks to hardware interrupts (kernel-level) or signals (user-level)
 - **Multiple threads (within the same process) can run in parallel on multiple CPUs**
- **Cooperative threads**
 - At most a single thread (within a given process) is allowed to run at a given point in time
 - A thread switch (within a given process) can only happen when:
 - The thread explicitly relinquishes the CPU (calls `yield()`)
 - The thread issues a blocking syscall (or terminates)


Preemptive vs. cooperative threads (continued)

- **Discussion**
 - Preemptive threads cause/expose more “race conditions” (i.e., concurrency bugs) because there are many more possible thread interleavings)
 - Cooperative threads provide a simpler programming model for concurrent tasks
 - Cooperative threads cannot take advantage of multiple CPUs
 - Cooperative threads may let a “misbehaving” thread monopolize the CPU ... but only up to the CPU share of the enclosing process
 - Before multiprocessor architectures became prevalent, many threading implementations were cooperative

Kernel threads vs. user threads

- “**Kernel threads**” (**kernel-managed threads**)
 - The kernel is aware that a process may encapsulate several schedulable execution contexts
 - The kernel manages these execution contexts
- “**User threads**” (**user-managed threads**)
 - Such execution contexts are managed from a library running in user level
 - the kernel is not aware of them, it only manages the encapsulating process, with a single execution context

Kernel threads




- **thread_create()** is implemented as a system call
- Faster than full process creation but still relatively heavy-weight

Limitations of kernel-level threads

- **Every thread operation must go through kernel**
 - Create, exit, join, synchronize or switch for any reason
 - On a modern processor, a syscall takes (approx.) 100+ cycles, while a function call takes 5 cycles
 - Result: threads 10x-30x slower when implemented in kernel
- One-size-fits-all thread implementation
 - Kernel threads must please all people
 - Maybe you pay for fancy features (priorities, etc.) that you do not need
- General heavy-weight memory requirements
 - E.g., requires a fixed-size stack within kernel
 - Other data structures designed for heavier processes

User threads



- An alternative: implement in user-level library
 - One kernel-thread per process
 - `thread_create()`, `thread_exit()`, ... are just library functions


Implementing user-level threads (as a library)

- Allocate a new stack for each invocation of `thread_create()`
- Keep a queue of runnable threads
- Replace some potentially blocking system calls (e.g., related to I/O: `read()`/`write()`/etc.)
 - If operation would block, switch and run different thread
- Schedule periodic timer signal (`setitimer()`)
 - Switch to another thread on timer signals (preemption)
- Multi-threaded web server example
 - Thread calls `read` to get data from remote client
 - “Fake” `read` wrapper function makes `read()` syscall in non-blocking mode – if no data available, schedule another thread
 - On timer tick or when idle, check which connections have new data, and switch to a thread that can make progress

Limitations of user-level threads

- **Cannot take advantage of multiple CPUs**
- **A blocking system call blocks all threads (within the same process)**
 - Some system calls can be replaced by non blocking ones (e.g., to read from network connections)
 - But, depending on the OS, this is not always possible for all potentially-blocking system calls (e.g., for disk I/O)
 - Such system calls may block all the threads of a given process
- **A page fault blocks all threads** (within the same process)
- Possible deadlock if one thread blocks on another
 - May block entire process and make no progress (more on deadlocks in another lecture)

Another possible threading design: user threads on (several) kernel threads



- User threads implemented on kernel threads
 - Multiple kernel-level threads per process
 - `thread_create()`, `thread_exit()` are still library functions
- **Sometimes called *N:M* threading (or *M:N*) or “hybrid” threading**
 - Have *N* user threads per *M* kernel threads
 - (“simple” user-level threads are *N:1* and “simple” kernel threads are *1:1*)

Limitations of $N:M$ threading

- Many of the same problems as $N:1$ threads
 - Blocked threads, deadlock, ...
- Hard to keep the number of kernel threads the same as available CPUs
 - The kernel knows how many CPUs are available
 - The kernel also knows which kernel-level threads are blocked
 - But tries to hide these things to applications for transparency
 - So a user-level thread scheduler might think that a thread is running while the underlying kernel thread is blocked
- The kernel does not know the relative importance of threads
 - Might preempt kernel thread in which library holds important lock

Advanced details

Threads: behavior upon `fork()`/`exec()`

- **What happens if one thread of a program calls `fork()`?**
 - Does the new process duplicate all threads? Or is the new process single-threaded?
 - Some Unix systems have chosen to have two versions of `fork()`
 - In general, only the calling thread is replicated in the child process
 - All of the other threads vanish in the child, without invoking thread-specific cleanup handlers
- **What happens if one thread of a program calls `exec()`?**
 - Generally, the program replaces the entire process, including all threads
 - Without invoking any thread-specific cleanup handler

Thread cancellation

- **One may want to cancel a thread before it has completed**
 - Example: when multiple threads concurrently search for a given data item in a database
 - Or when you hit the stop button of a Web browser, all the threads in charge of loading the code of the web page and the various images should be cancelled
- **Asynchronous cancellation**
 - One thread immediately terminates the target thread
 - Main issue: what if resources have been allocated and/or the target thread is in the midst of updating data shared with other threads?
 - May lead to incoherent state
- **Deferred cancellation**
 - The target thread periodically checks whether it should terminate, giving it an opportunity to terminate itself in an orderly fashion
 - Such points are called cancellation points

Signal handling

- Handling signals in a single-threaded program is straightforward
- In a multi-threaded program, who should receive the signal? Several possibilities:
 - Deliver the signal to the thread to which the signal applies (e.g., `SIGSEGV`)
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process
- POSIX threads have the `pthread_kill(pthread_t tid, int signal)` function
- In many Unix systems, the decision is usually made as follows:
 - Only a single thread receives a given signal instance within a process
 - If the signal is clearly related to a given thread, select this one
 - E.g., in case of a hardware fault (like `SIGSEGV`), or a call to `pthread_kill()`
 - Otherwise, select an arbitrary thread within the process

Thread-specific data

- **All threads share the data of the enclosing process**
- In some circumstances, each thread may need to have its own copy of certain data
- Most thread libraries provide some support for thread-specific data:
 - POSIX Thread-specific data (a relatively complex API)
 - “Thread local storage” (non-standard but simpler and implemented in different Unix variants like Linux, FreeBSD and Solaris)
- Thread-local storage – example:
 - Simply include the **`__thread`** specifier in the declaration of a global or static variable
 - Example: **`static __thread char buf[BUF_SIZE];`**

Thread pools

- A Web server could create a thread to handle each client request
 - Although it is cheaper than creating a process, creating a thread is costly, especially regarding the request service time
 - If there is no bound on the number of concurrently active threads, we could exhaust the system resources (CPU, RAM) and cause thrashing
- Thread pools address these two issues
 - Create a number of threads when the (server application) process starts and place them into a pool where they wait for work
 - When a server receives a request, it awakens a thread from the pool if any available and waits otherwise
 - When the thread has finished servicing the request, it returns to the pool, awaiting for more work